

Advanced Machine Learning - Exercise 4

May 16, 2018

1. In this exercise you will implement a variational autoencoder using Keras, a powerful high-level library that works on top of TensorFlow (and a few additional libraries).
 - (a) (Do not submit) Get familiar with Keras, <https://keras.io/getting-started/sequential-model-guide/>
 - (b) (Do not submit) Download and go over the python code, `variational_autoencoder.py`, available on the course Moodle.
 - (c) The code provided defines and trains a VAE. Add an encoder which maps MNIST digits to the latent space. Using this encoder, visualize the test set in the latent space (make sure to use Keras functional API when defining the encoder, <https://keras.io/getting-started/functional-api-guide/>). Take one image per digit and print its corresponding mapping coordinates in the latent space, present the answer as a table.
 - (d) The VAE is a generative model, we would like to be able to generate new MNIST digits based on the VAE we have trained. Use the following code to define a generator that based on a sample from the latent space, generates a digits.

```
decoder_input = Input(shape=(latent_dim,))
_h_decoded = decoder_h(decoder_input)
_x_decoded_mean = decoder_mean(_h_decoded)
generator = Model(inputs=decoder_input, outputs= _x_decoded_mean)
```

Given a sample from the latent space, for instance, `z_sample = np.array([[0.5, 0.2]])`. Applying the generator, `x_decoded = generator.predict(z_sample)`, results with a generated digit.

- (e) Add a sequence of images which form an interpolation from one digit to another. Namely
 - Take two original images from MNIST of different digits.
 - Find their representation in the latent space.

- Sample 10 points from the line connecting the two representations in the latent space and generate their images using the generator.
- (f) The code provided models both the mean and the variance. Fix the variance vector to constant ones and train again when learning only the mean vector. Answer sections (c), (d) and (e) again. (note that in the code the variance variable holds the log of the values, therefore you should set the values to zero).
- (g) The autoencoder architecture provided is a simple fully connected neural network. Change `intermediate_dim` parameter in the code to 16 and change the architecture to a ConvNet with the following structure:
- Conv2D(16, (3, 3), activation='relu', padding='same')
 - MaxPooling2D((2, 2), padding='same')
 - Conv2D(8, (3, 3), activation='relu', padding='same')
 - MaxPooling2D((2, 2), padding='same')

The corresponding decoder which you need to define as well will have the following structure:

- Conv2D(8, (3, 3), activation='relu', padding='same')
- UpSampling2D((2, 2))
- Conv2D(16, (3, 3), activation='relu')(x)
- UpSampling2D((2, 2))(x)
- Conv2D(1, (3, 3), activation='sigmoid', padding='same')

Note that the ConvNet is to be added in addition to the fully connected layers. Namely, your network should still output `z_mean` and `z_log_var` with a dimension of `latent_dim` each (see `conv_arch.py` on the course Moodle, the code will require additional changes to work properly). Train the VAE with the new architecture and answer sections (c), (d) and (e) again.

Submission Guidelines:

- Add your IDs and the path to your code's folder in the following Google Form: <https://goo.gl/YmRMrc>
 - No need to print code.
 - Make sure your code's folder has the proper read permissions.
 - Your code should be readable and well-documented.
2. Let P, Q be two normally distributed random variables, $P \sim \mathcal{N}(\mu_1, \sigma_1^2)$ and $Q \sim \mathcal{N}(\mu_2, \sigma_2^2)$. Derive an expression for the Kullback–Leibler divergence of P and Q , $D_{KL}(P||Q)$.